

QSS20: Modern Statistical Computing

Session 04: Merging and basic regex

Goal for next few sessions

- ▶ Some course housekeeping
- ▶ Exact matching: types of joins
 - ▶ Inner joins
 - ▶ Outer joins
 - ▶ Left joins
 - ▶ Right joins
- ▶ Basic regex for two purposes:
 1. Clean join fields for exact matching/merges
 2. Clean join fields for fuzzy/probabilistic matching/merges
- ▶ Fuzzy/probabilistic matching and merges

Goal for next few sessions

- ▶ **Some course housekeeping**
- ▶ Exact matching: types of joins
 - ▶ Inner joins
 - ▶ Outer joins
 - ▶ Left joins
 - ▶ Right joins
- ▶ Basic regex for two purposes:
 1. Clean join fields for exact matching/merges
 2. Clean join fields for fuzzy/probabilistic matching/merges
- ▶ Fuzzy/probabilistic matching and merges

PSET 1 Submission

- ▶ **Group:** tonight 1159 PM EST (will give extra time during break for u guys to meet and coordinate submission)
- ▶ **Individual:** same deadline unless using some/all of four free late days. If using all four, due at Saturday 04.25 at 1159 PM EST
- ▶ I'll give an emoji reaction on issue when i've seen it so you know it worked; i'll also comment on issue if i'm having trouble running your code

Organization of activity-based practice code

https://github.com/rebeccajohnson88/qss20_slides_activities#readme

These are jupyter notebook-based activities to practice Python or other concepts:

- [00_latex_output_examples_solutions.ipynb](#)
 - Data: DC crime reports in 2020
 - Concepts covered:
 - Writing a pandas dataframe or table to use in LaTeX
 - Row filtering
 - Saving figures
 - Iterating and saving figures with informative names
- [01_pandas_datacleaning_examples.ipynb](#)
 - Data: sample of Chicago health/hygiene inspection results
 - Concepts covered:
 - Cleaning column names (eg subbing out spaces and changing to lowercase)
 - Checking datatypes within a pandas dataframe and recasting
 - Creating new true/false variables using `np.where`
 - Creating new categorical variables that involve recoding an existing categorical variable using `map` and a dictionary
- [02_more_pandas_datacleaning.ipynb](#)
 - Data: DC crime reports in 2020
 - Concepts covered:
 - Aggregation using `groupby` and `agg`
 - Lambda functions within aggregation
 - Recoding variables using `np.where`
 - Recoding variables using `np.select`
 - Recoding variables using `map` and dictionary
 - Loop to find matches within a broader pool of data
 - Function to find matches within a broader pool of data

Updated course schedule

https://rebeccajohnson88.github.io/qss20/docs/course_schedule.html

Tuesday 04-20	Intro to merging	Problem set one
Thursday 04-22	Merging: probabilistic merge and more regex	
Tuesday 04-27	Merging (continued) and PSET 1 review	Regular expressions for pattern matching Final project step 1
Thursday 04-29	SQL via Python	
Tuesday 05-04	Text as data part one	Final project step 2
Thursday 05-06	Text as data part two	
Tuesday 05-11	TBD	
Thursday 05-13	Python: spatial data using geopandas	Problem set two
Tuesday 05-18	Python: reading data from APIs and basic web scraping	
Thursday 05-20	High-performance computing	Final project step 3
Tuesday 05-25	TBD	
Thursday 05-27	Workflow: Beamer and Tikz graphics	
		Slides for final

Steps towards final project

1. Make sure you can access this private repo and DM me if you need re-sent invite:
https://github.com/rebeccajohnson88/qss20_s21_proj
2. Join #sip_finalproject on Slack
3. Will post details on Canvas tomorrow for Final project Step 1, due Tuesday 04.27 alongside the DataCamp assignment, but broadly: (1) sign up for background reading, (2) copy over LaTeX/Overleaf template I'll share, and (3) write < 1 page memo outlining data used in the background reading, key takeaways, and interesting and feasible follow-up questions

Mid-term evaluation of our course

- ▶ Will circulate anonymous feedback survey later this week covering course pace, clarity, and what's going more versus less well

Goal for next few sessions

- ▶ Some course housekeeping
- ▶ **Exact matching: types of joins**
 - ▶ Inner joins
 - ▶ Outer joins
 - ▶ Left joins
 - ▶ Right joins
- ▶ Basic regex for two purposes:
 1. Clean join fields for exact matching/merges
 2. Clean join fields for fuzzy/probabilistic matching/merges
- ▶ Fuzzy/probabilistic matching and merges

Working example: have dataset on Dartmouth students and want to merge in background information about their district

► **Main or “left” dataset**

Student	Year	District	NCES ID
Rebecca	2021	New Trier High School	1728200
Jennifer	2022	Hanover High	3302670
Jason	2022	Homeschool	NA
⋮			


► **Auxiliary or “right” dataset**

District	NCES ID	% FRPL
New Trier HS	1728200	X%
Hanover HS	3302670	Y%
Lebanon HS	4107380	Z%
⋮		


Possible join keys

- ▶ **Unique identifier:** used for “exact matching” — or a Yes/No match on that basis
 - ▶ E.g., is the NCES ID of New Trier found in the dataset of demographics?
- ▶ **Other identifiers:** can be used for either “exact match” or for “probabilistic/fuzzy matching”
 - ▶ **Probabilistic:** what’s the likelihood that “New Trier district” and “New Trier HS” are the same entity?


Conceptual overview of four types of joins




Left Join



Right Join



Inner Join



**Full Outer
Join**

Source: Trifacta

Inner join in this context

In words: “drop all students whose districts don’t appear in the demographics data; drop all districts that don’t appear in the Dartmouth student data”

► **Main or “left” dataset**

Student	Year	District	NCES ID
Rebecca	2021	New Trier High School	1728200
Jennifer	2022	Hanover High	3302670
Jason	2022	Homeschool	NA
⋮			

► **Auxiliary or “right” dataset**

District	NCES ID	% FRPL
New Trier HS	1728200	X%
Hanover HS	3302670	Y%
Lebanon HS	4107380	Z%
⋮		

Outer join in this context

In words: “keep all students from the student-level data; keep all schools from the school-level data; even if there’s not an overlap”

Student	Year	District	NCES ID	% FRPL
Rebecca	2021	New Trier High School	1728200	X%
Jennifer	2022	Hanover High	3302670	Y%
Jason	2022	Homeschool	NA	NA
NA	NA	NA	4107380	Z%
⋮				

Left join in this context

In words: “keep all students from the student-level data; drop any school from the school-level data that doesn’t merge onto a student”

► **Main or “left” dataset**

Student	Year	District	NCES ID
Rebecca	2021	New Trier High School	1728200
Jennifer	2022	Hanover High	3302670
Jason	2022	Homeschool	NA
⋮			

► **Auxiliary or “right” dataset**

District	NCES ID	% FRPL
New Trier HS	1728200	X%
Hanover HS	3302670	Y%
Lebanon HS	4107380	Z%
⋮		

Right join in this context

In words: “drop students who don’t have a school in the school-level data; keep all schools from the student-level data even those that don’t merge onto any student”

► **Main or “left” dataset**

Student	Year	District	NCES ID
Rebecca	2021	New Trier High School	1728200
Jennifer	2022	Hanover High	3302670
Jason	2022	Homeschool	NA
⋮			

► **Auxiliary or “right” dataset**

District	NCES ID	% FRPL
New Trier HS	1728200	X%
Hanover HS	3302670	Y%
Lebanon HS	4107380	Z%
⋮		

How do we code these different types of joins in practice?
Example with left join and join key has same colname in both

```
1  
2 ## perform a left join on the student data  
3 ## and schools data  
4 stud_wschoool = pd.merge(students ,  
5                           schools ,  
6                           how = "left" ,  
7                           on = "NCES ID" ,  
8                           indicator = "student_mergestatus" )
```

- ▶ **how**: argument to tell it inner, left, right, outer, or cross; defaults to inner
- ▶ **on**: name of join key (in this case single key)
- ▶ **indicator**: optional arg to add a col to the resulting data (string is what to call it) that helps diagnose merge status; good for post-merge dx

Example with inner join and join key has different name

```
1  
2 ## perform a left join on the student data  
3 ## and schools data  
4 stud_wschoool = pd.merge(students ,  
5                         schools ,  
6                         how = "inner" ,  
7                         left_on = "NCES ID" ,  
8                         right_on = "ncesnumeric")
```

Example with left join and multiple join keys

```
1  
2 ### perform a left join on the student data  
3 ### and schools data  
4 stud_wschoool = pd.merge(students ,  
5                           schools ,  
6                           how = "left" ,  
7                           left_on = ["NCES ID" ,  
8                                       "Dist name" ] ,  
9                           right_on = ["ncesnumeric" ,  
10                                      "distnamechar" ] ,  
11                          indicator = "student_mergestatus" )
```

Non-exhaustive checklist of merge diagnostics

1. How many rows were in each data before the merge? What about after?
2. If doing a left join, did we properly retain all left-hand side rows?
3. **For strings as join keys:** if a lot of rows were lost in a merge, could that be due to spelling/punctuation variations in a character join key?
4. **For numeric identifiers as join keys:** if a lot of rows were lost in a merge, could that be due to things like the id having leading zeros and those being stripped at some stage? (e.g., one dataset identifies an entity as 002548; another as 2548)

Next up: basic regex to improve match rates for strings as join keys

- In example below, what if we didn't have the NCES ID numeric identifier? Ways to improve match rates for spelling variations (sometimes called `entity resolution`)

Student	Year	District
Rebecca	2021	New Trier High School
Jennifer	2022	Hanover High
Jason	2022	Homeschool
⋮		

District	% FRPL
New Trier HS	X%
Hanover HS	Y%
Lebanon HS	Z%
⋮	