

# QSS20: Modern Statistical Computing

## Session 06: Basic regex

# Goal for next few sessions

- ▶ Exact matching: types of joins
  - ▶ Inner joins
  - ▶ Outer joins
  - ▶ Left joins
  - ▶ Right joins
- ▶ **Basic regex for two purposes:**
  1. Clean join fields for exact matching/merges
  2. Clean join fields for fuzzy/probabilistic matching/merges
- ▶ Fuzzy/probabilistic matching and merges

# Today: basic regex to improve match rates for strings as join keys

- In example below, what if we didn't have the NCES ID numeric identifier? Ways to improve match rates for spelling variations (sometimes called `entity resolution`)

<b>Student</b>	<b>Year</b>	<b>District</b>
Rebecca	2021	New Trier High School
Jennifer	2022	Hanover High
Jason	2022	Homeschool
⋮		

<b>District</b>	<b>% FRPL</b>
New Trier HS	X%
Hanover HS	Y%
Lebanon HS	Z%
⋮	

## Working example

Want to clean school names and classify them into different types (elementary; middle school; high school; charter; alternative; etc...). E.g.:

Central Columbia Ms

Riley County High

Jarrell H S

Trumbull School

SAN GABRIEL ELEMENTARY

Plains El

Pond Hill School

Franklin Elementary

P.S. 119

ANDREW CARNEGIE MIDDLE

Example of variety of names that all match the pattern within `str.contains`

```
1 cep_optin['is_elem'] =  
2 np.where(cep_optin.schoolname_lower.str.contains("\s+elem",  
3 regex = True), True, False)
```

Examples of True show a lot of variation that could make merges to other data difficult...

paint branch elementary

stewart county elementary school

stove prairie elementary school

winchester avenue elementary school

oak hill elem.

lewis and clark elem.

saunemin elem school

desert springs elementary school

fifth district elementary

linden elementary school

## re module provides more flexible alternative to pandas str methods

- ▶ Need to import at top: `import re`
- ▶ General structure: `re.something(r'‘somepattern’', some_str)`
- ▶ Challenging part is constructing the pattern that captures what you want to capture

## First example

- We want to standardize the school names so that if it's an elementary school, it has the string “elemschool” after its other identifiers. E.g.:

original	cleaned
paint branch elementary	paint branch elemschool
stewart county elementary school	stewart county elemschool
stove prairie elementary school	:
winchester avenue elementary school	
oak hill elem.	
lewis and clark elem.	
saunemin elem school	
desert springs elementary school	
fifth district elementary	
linden elementary school	

## Approach 1 using re.sub

Construct a pattern to match all identified variations you want to sub out. One natural inclination is to try to enumerate the variations and use | (or), e.g.:

```
1 ## define pattern
2 elem_pattern = r"elementary|elem|elem\\.|elementary school"
3
4 ## replace in one string
5 one_str_clean = re.sub(elem_pattern,
6                        "elemschool", one_str)
7
8 ## replace for all strings in the column schoolname_lower
9 all_str_clean = [re.sub(elem_pattern, "elemschool",
10                        one_str)
11                  for one_str in df.schoolname_lower]
```

But this leads to issues...



## Issues with earlier pattern

(1) Sees elementary and subs it out but leaves school in; (2) leaves in . after elem. since matches elem

orig_name	cleaned_name
paint branch elementary	paint branch elemschool
stewart county elementary school	stewart county elemschool school
stove prairie elementary school	stove prairie elemschool school
winchester avenue elementary school	winchester avenue elemschool school
oak hill elem.	oak hill elemschool.
lewis and clark elem.	lewis and clark elemschool.
saunemin elem school	saunemin elemschool school
desert springs elementary school	desert springs elemschool school
fifth district elementary	fifth district elemschool
linden elementary school	linden elemschool school

# Making the pattern more robust to variations

```
1 ## define pattern  
2 elem_pattern_try2 = r"(elem.*)(\s+)?(school)?"
```

Breaking down each component:

- ▶ Creating groups using ( ): these help us define groups of characters to look at together
- ▶ elem.\*: matches elem, elem., and elementary (would maybe want to make more restrictive if we had schools named things like element that we didn't want to match)
- ▶ Multiple spaces: uses “metacharacter” to match 1+ spaces:  
    \s+
- ▶ ? : **optional pattern**, or match if this pattern occurs but also match if it doesn't (in this case, it's useful since schools like fifth district elementary have nothing afterwards)
- ▶ (school)? : similarly, sometimes ends with school and we want to replace; other times just ends with elementary or elem

# Executing with re.sub

```
1 ## define pattern
2 elem_pattern_try2 = r"(elem.*)(\s+)?(school)?"
3
4 ## replace for all strings in the column schoolname_lower
5 all_str_clean_try2 = [re.sub(elem_pattern_try2, "elemschool",
6                             one_str)
7                        for one_str in df.schoolname_lower]
```

## Better result than our first attempt :)

<b>orig_name</b>	<b>cleaned_name</b>
paint branch elementary	paint branch elemschool
stewart county elementary school	stewart county elemschool
stove prairie elementary school	stove prairie elemschool
winchester avenue elementary school	winchester avenue elemschool
oak hill elem.	oak hill elemschool
lewis and clark elem.	lewis and clark elemschool
saunemin elem school	saunemin elemschool
desert springs elementary school	desert springs elemschool
fifth district elementary	fifth district elemschool
linden elementary school	linden elemschool

# Other re operations

- ▶ In previous example, we:
  1. Defined a pattern: different variations of elementary school
  2. Used `re.sub(pattern, replacement, string)` to substitute something else when we match the pattern
- ▶ In other examples, we may want to:
  - ▶ Define a pattern that characterizes different subparts of a string (e.g., elementary is one part we want to extract; school is another)
  - ▶ Match that pattern
  - ▶ Extract the matches and do something

First, we can use “metacharacters” or shortcuts to match general types of patterns

<https://www.geeksforgeeks.org/python-regex-metacharacters/>

Common ones:

- ▶ `\d` or `[0-9]`: numbers
- ▶ `[A-Z]`: uppercase alpha (any; can do subsets like `[ABC]`)
- ▶ `[a-z]`: lowercase alpha
- ▶ `\w`: alphanumeric (so numbers of letters)
- ▶ `+`: match one or more appearances (e.g., if we have two usernames—`rebeccajohnson8`; `rebeccajohnson88`; `rebeccajohnson796`—could do `"[a-z]+\d+"` to match all versions)
- ▶ `*`: match anything
- ▶ `^`: match at beginning
- ▶ `$`: match at end
- ▶ `{x, y}`: match a pattern of length `x` to length `y` (e.g., if rather than capture something repeated 1 or more times, we wanted to capture numbers that look like ages and could be length 1-3, write as `"\d{1,3}"`)

## Then, three methods that do similar things

1. `re.findall(pattern, string)`
  - ▶ **What it does/returns:** scans the string from left to right and returns the matches in a **list**
  - ▶ **Useful for:** parsing a string and then recombining elements (e.g., elementary could be list element 0; school could be list element 1)
2. `re.search(pattern, string)`
  - ▶ **What it does/returns:** scans the **entire string** and returns the substring(s) regardless of where it appears in the string. If no matches found, returns `None`. If matches found, returns **MatchObject**
  - ▶ **Useful for:** checking *if* there's a match (since can check whether result is equal to `None`); same use above of getting substrings
3. `re.match(pattern, string)`
  - ▶ **What it does/returns:** Operates similarly to `re.search()` but rather than searching the entire string, only returns if found at beginning of string

Both `re.match()` and `re.search()` only return the first appearance of a pattern; if want to return all occurrences (e.g., count how many times “Trump” appears in a single tweet), can use either `re.findall()` or `re.finditer()`

## Focusing on re.findall: executing

**Example:** match words before and after the charter pattern

```
1 ## charter pattern
2 charter_pattern = r"(.*)\s+(charter)(\s+)?(\w+)?"
3
4 ## findall
5 test_charter_findall = [re.findall(charter_pattern ,
6                                   school)
7                           for school in combined_examples]
```



# What this returns: a list of lists

## orig\_name

buffalo collegiate charter school  
thomas edison charter academy  
moving everest charter school  
life source international charter  
south valley academy charter school  
neighborhood charter school of harle  
brighter choice charter school-girls  
children's community charter  
frontier elementary school  
columbus humanities, arts and...  
okemos public montessori-central  
pawhuska es  
east valley senior high  
glenpool es  
number 27  
south fork elementary

```
: [[('buffalo collegiate', 'charter', ' ', 'school')],  
  [('thomas edison', 'charter', ' ', 'academy')],  
  [('moving everest', 'charter', ' ', 'school')],  
  [('life source international', 'charter', ' ', 'school')],  
  [('south valley academy', 'charter', ' ', 'school')],  
  [('neighborhood', 'charter', ' ', 'school')],  
  [('brighter choice', 'charter', ' ', 'school')],  
  [('children's community', 'charter', ' ', 'school')],  
  [],  
  [],  
  [],  
  [],  
  [],  
  [],  
  [],  
  [],  
  []]
```

What if we use same pattern and use `re.search` instead on the list of school names? How to execute

```
1 ## charter pattern
2 charter_pattern = r"(.*)\s+(charter)(\s+)?(\w+)"
3
4 ## search
5 test_charter_search = [re.search(charter_pattern ,
6                               school)
7                        for school in combined_examples]
```



MatchObject isn't useful in its own right; how do we extract matches? `.group()` method on MatchObject

Example with: thomas edison charter academy

```
1 ##### here, we're just focusing on the 2nd match (thomas edison
   charter academy)
2 ##### and we're getting the first group from that match
3 thomas_match = test_charter_search[1]
4 thomas_match
5
6 ##### example where we're just getting the first group
7 ##### (name of school before charter)
8 thomas_firstgroup = thomas_match.group(1)
```

Prints: thomas edison

# Generalizing to show all groups for that one match

```
### iterate over all groups and print
for i in range(0, len(thomas_match.groups())+1):
    print("Group " + str(i) + " is: ")
    print(thomas_match.group(i))
|
## see error if we go beyond actual number of
## groups thomas_match.group(5)
```

```
Group 0 is:
thomas edison charter academy
Group 1 is:
thomas edison
Group 2 is:
charter
Group 3 is:

Group 4 is:
academy
```

# Break for activity



Activity section at end of [04\\_basicregex\\_blank.ipynb](#)